

Deep Learning Practical Work 1-a and 1-b

Images descriptors with SIFT and Bag of Words

Nicolas Thome

Guillaume Couairon

Mustafa Shukor

Asya Grechka

Matthieu Cord

Homework

Data, code, and PDF version of this file are available at
<https://rdfsia.github.io>

Goals

During these first practical works, we are going to prepare the development of an initial image classification model. The goal is to produce an algorithm that is able to classify images from the dataset 15-Scenes. This dataset contains 4485 images belonging to 15 scenes categories (e.g. kitchen, bedroom, street, etc.)

First, in 1-a, we are going to use **SIFT ((Scale-invariant feature transform))** which are local visual descriptors and thus can encode a small *patch* of image. Then, still in 1-a, we'll explore some archetype descriptors that represent recurring patterns to build a **visual dictionary**. Then, we'll aggregate, in 1-b, the SIFT of each image with a **BoW (Bag of Words)** method that express in a compact way an image.

Finally, in 1-c, we'll learn to classify each image, from its BoW representation by using a **SVM (Support Vector Machine)**.

Data · Code · Python

During our practical work sessions, we'll use Python3 alongside very popular *packages* in scientific computing and machine learning (numpy, matplotlib, scikit-learn, pytorch, ...). A quick introduction of Python is proposed here <https://deep-learning-polytech.github.io/python.html>.

All the coding sessions will be done on Google colab, but you can still download a jupyter notebook from it if you prefer running the code on your local machine. Be aware that we won't help you if you have a problem related to your machine in that case.

Section 1 – SIFT

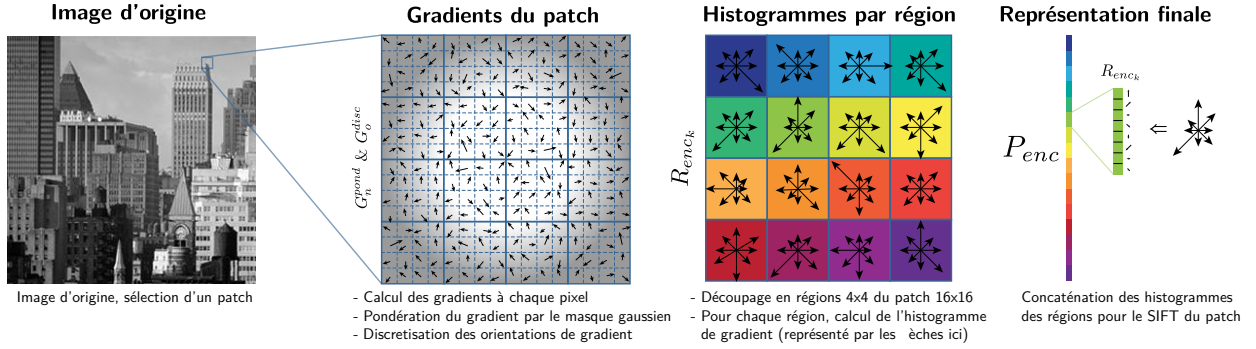


Figure 1: Illustration of the SIFT method applied to a patch.

A SIFT (Scale-invariant feature transform) is a visual and local descriptor. It can transform a small image *patch* of size 16×16 pixels into a numerical representation (a vector of 128 dimension in that case) which will be robust to small perturbations and transformations. Meaning that two similar patches will have similar SIFTs.

1.1 Computing the gradient of an image

First, we want to compute the gradient of an image at a pixel (x, y) :

$$G(x, y) = \left[\frac{\partial I}{\partial x} \quad \frac{\partial I}{\partial y} \right]^\top = [I_x \quad I_y]^\top \quad (1)$$

In practice, the partial derivatives will be approximated by finite differences: in that case, the partial derivatives I_x et I_y can be obtained using the convolution (denoted by \star) of the image by a kernel $I_x = I \star M_x$, $I_y = I \star M_y$. We'll use the following 3×3 Sobel kernels:

$$M_x = \frac{1}{4} \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} \quad M_y = \frac{1}{4} \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix} \quad (2)$$

From I_x and I_y , we can compute the norm of the gradient $G_n = \|G\| = \sqrt{I_x^2 + I_y^2}$ and the direction G_o (code is given for that).

Questions

1. Show that kernels M_x and M_y are separable, i.e. that they can be written $M_x = h_y h_x^\top$ and $M_y = h_x h_y^\top$ with h_x and h_y two vectors of size 3 to determine.
2. Why is it useful to separate this convolution kernel?

1.2 Computing the SIFT representation of a patch

Let's define the algorithm that can produce the SIFT representation of an image patch P of 16×16 pixels. The whole method that we will cover is illustrated in the figure 1.

The patch P is going to be decomposed in 16 sub-regions R_i of 4×4 pixels each. Each region will be encoded as a vector $R_{enc_k} \in \mathbb{R}^8$. To do so, we'll use the SIFT descriptor (Scale Invariant Feature Transform) which consists in computing a histogram of gradient directions.

The complete patch will be described by a vector $P_{enc} \in \mathbb{R}^{128}$ which will be the concatenation of all regions encoding.

Algorithm

The SIFT algorithm is as follow:

- For each pixel of the patch P , we compute the image gradients with respect to each pixel will be stored as a gradient direction matrix $G_o \in \mathbb{R}^{16 \times 16}$ and a gradient norm matrix $G_n \in \mathbb{R}^{16 \times 16}$.
- We discretize the gradient direction matrix into 8 bins / values. The value 0 will be associated to the direction 0° (north), the value 1 to the direction 45° (north-east), the value 2 to the direction 90° (east), etc. Thus, we obtain a matrix $G_o^{disc} \in \mathbb{N}^{16 \times 16}$.
- We compute a 2d gaussian $N \in \mathbb{R}^{16 \times 16}$ centered on the center of P and with standard deviation $\sigma = 0.5 \times width$ with $width = 16$ pixels. We weight G_n by N to obtain G_n^{pond} .
- For each region of size 4×4 , we compute a histogram of gradient directions $R_{enc_k} \in \mathbb{R}^8$. Each bin of the histogram is the sum of the gradient weighted norms for a given direction. For example, $R_{enc_k}[0]$ is the sum of values of $G_n^{pond}[i, j]$ such as (i, j) is in the region k and that $G_o^{disc}[i, j] = 0$.
- We concatenate our 16 vectors R_{enc_k} to make the descriptor $P_{enc} \in \mathbb{R}^{128}$.

We conclude the algorithm by a **post-processing** on P_{enc} :

- We compute the L_2 norm of the descriptor P_{enc} . If it is smaller than a threshold of 0.5, P_{enc} is set to a null vector and is immediately returned.
- Otherwise, we normalize the descriptor to have a unit euclidean (L_2) norm ($\|P_{enc}\|_2 = 1$).
- Finally, values greater than 0.2 will be clamped to 0.2, and then the descriptor is L_2 normalized a second time.

Practice

In the beginning of the colab, a file `tools.py` is downloaded which contains many useful functions that you can use. Write your code where indicated in the colab notebook.

- Write the function `compute_grad(I)` which computes the gradient of the input image and returns I_x and I_y . You'll use the function `conv_separable(I, ha, hb)` which computes $I \star M$ with $M = h_a \times h_b^\top$.
- Write the function `compute_grad_mod_ori(I)` which return G_n and the discretized orientations G_o^{disc} of the input image. You'll use `compute_grad(im)` and `compute_grad_ori(Ix, Iy, Gn)`.

- From now on, we don't work anymore on the whole image but rather on a patch of size 16×16 . Read the function `compute_sift_region(Gn, Go, mask=None)` which returns the representation P_{enc} of a patch given G_n and G_o^{disc} and also an optional gaussian mask which can help produce G_n^{pond} if provided. Note that this function will contain the post-processing of SIFT. Try to understand this function and match it to the algorithm given in this PDF.

The gaussian mask to give in input is generated by the function `gaussian_mask()` which is provided.

- The function `compute_sift_region` calls the function `compute_histogram(g_n, g_o)` that you have to fill.

Test your function manually and visually thanks to the function

`display_sift_region(I, compute_grad_mod_ori, compute_sift_region, x, y)` which takes an image in input, coordinates x, y where to cut a patch, and your computing functions.

Questions

3. What is the goal of the weighting by gaussian mask?
4. ★ Explain the role of the discretization of the directions.
5. Justify the interest of using the different post-processing steps.
6. ★ Explain why SIFT is a reasonable method to describe a patch of image when doing image analysis.
7. ★ ★ Interpret the results you got in this section.

1.3 Computing SIFTs on the image dataset

To compute the SIFT representation of an image, we must choose center points whose surrounding will be used to compute the SIFTs. Multiple methods¹ exist to efficiently choose them but a simple way is to densely sample them. i.e. to take a patch every n pixels. In our case, we will sample one 16×16 patch every 8 pixels. Thus note that the sampled patches will have some overlap!

- Write the function `compute_sift_image(I)` which computes the SIFTs of an image.

The function `x, y = dense_sampling(I)` provides two lists `x` and `y`. The upper-left coordinates of every patches are the pairs (x_i, y_i) . The SIFTs are stored in a numpy array of size $n_x \times n_y \times d_{SIFT}$ with n_x the number of coordinates alongside the x -axis and n_y alongside the y -axis, and d_{SIFT} the size of a SIFT.

Section 2 – Visual Dictionary

Theory

We are going to use all SIFTs descriptors P_{enc} extracted from our dataset to compute a **visual dictionary**. A *word* in this dictionary is somekind of SIFT-like descriptor likewise $c_m \in \mathbb{R}^{128}$. The goal of this dictionary is to represent as fidely as possible all the SIFTs but using a limited amount of words. Thus, those words will be frequent patterns in the extracted SIFTs.

Concretely, our goal is to determine K centers (clusters) c_m of the SIFT space that minimize a distance from the actual SIFTs x_i . This is our objective:

¹for the curious see the *Difference of Gaussian* at https://docs.opencv.org/master/da/df5/tutorial_py_sift_intro.html

$$\min_{c_m} \sum_{m=1}^M \sum_{x_i \in C_m} \|x_i - c_m\|_2^2 \quad (3)$$

A classic, yet efficient, way to solve this problem is to use the **K-Means** algorithm whose behavior is rather simple:

- Initialize M centers c_m from randomly-sampled points x_i .
- Alternate until convergence:
 - Assign to each point x_i the closest center
 - Re-compute the centers c_m as the average of their assigned points

In Practice

With the function `compute_load_sift_dataset`, we can obtain a list where each element is a list of the SIFTs of a single image.

You must complete the function `compute_visual_dict` so that it computes the visual dictionary, i.e. the M clusters. To do so, we are going to use scikit-learn : <http://scikit-learn.org/stable/modules/clustering.html#k-means>

Remember to add a dummy vector filled with zeros as “cluster” which will represent the null SIFTs.

Choose (but you can try other values) 1000 clusters to determine with K-Means. At the end, you should have 1001 clusters by adding the zero cluster.

You can then use the function `compute_or_load_vdict` to compute the clusters on all the SIFTs.

The function `get_regions_and_sifts` will get you the patches of an image alongside the corresponding SIFTs. To analyse the clusters of a visual dictionary, we could search for the image patches that are their closest (in the SIFT space!). Try different things to get a deeper understanding of your visual dictionary. For example, you could look for the 50 patches closest to a particular cluster, or for each cluster display the closest patch, etc. To help you, the function `display_images` can display regions.

Questions

8. ★ Justify the need of a visual dictionary for our goal of image recognition that we are currently building.
9. Considering the points $\{x_i\}_{i=1..n}$ assigned to a cluster c , show that the cluster's center that minimize the dispersion is the barycenter (mean) of the points x_i :

$$\min_c \sum_i \|x_i - c\|_2^2 \quad (4)$$

10. In practice, how to choose the “optimal” number of clusters?
11. Why do we create a visual dictionary from the SIFTs and not directly on the patches of raw image pixels?
12. Comment the results you get.

Section 3 – Bag of Words (BoW)

Theory

In this section, the goal is to obtain a numerical representation of each image that will help us later to classify this image (e.g. tell what does the image represent).

For now, we have for each image an ensemble of local SIFT descriptors (not necessarily the same amount per image), and a visual dictionary made of “mean” descriptors.

The goal of BoW is to summarize all the local descriptors of an image into a global descriptor with the help of the visual dictionary.

For an image I , we have all its descriptors x_i in a matrix $X \in \mathbb{R}^{n_{patch} \times d}$. We also have the visual dictionary under a matrix form $C \in \mathbb{R}^{M \times d}$.

The creation of the BoW descriptor of the image I is done in two steps:

- The **encoding** step: each descriptors x_i is encoded as a vector $h_i \in \mathbb{R}^M$ by using the visual dictionary. We get a new matrix $H \in \mathbb{R}^{n_{patch} \times M}$.
In our case, we will use a “nearest-neighbors” encoding: the vector $h_i \in \mathbb{R}^M$ is a one-hot vector indicating which word of the dictionary is the closest:

$$h_i[j] = \begin{cases} 1 & \text{if } j = \arg \min_k \|x_i - c_k\|^2 \\ 0 & \text{otherwise} \end{cases} \quad (5)$$

- The *pooling* step: we aggregate the h_i over all local descriptors to obtain a vector $z \in \mathbb{R}^M$ that globally describe the image.
In our case, z will simply be the sum of h_i .

The descriptors will finally be normalized with a euclidean norm (L_2).

Pratique

Write a function `compute_feats` that takes in input the visual dictionary matrix C and the matrix X containing all SIFTs of an image. This function will return the BoW representation z of the image.

Use the code given in the notebook to visualize the BoW representation.

Questions

- ★ Concretely, what does the vector z represent of the image?
- ★ ★ Show and discuss the visual results you got.
- What is the interest of the nearest-neighbors encoding? What other encoding could we use (and why)?
- What the interest of the sum pooling? What other pooling could we use (and why)?
- What is the interest of the L_2 normalization? What other normalization could we use (and why)?